

## GERÊNCIA DE INFRAESTRUTURA PARA BIG DATA

Prof. Tiago Ferreto – tiago.ferreto@puccs.br



## HADOOP

### Big Data and Hadoop History

- Early 2000s – emergence of storage and processing methodologies (“Big Data”) from search engine providers – principally Google and Yahoo!
- Search engine providers
  - First group to face with Internet scale problems – how to process and store indexes of all of the documents in the Internet
- Timeline
  - 2003 – Google releases “The Google File System” whitepaper – scalable and fault tolerant file system
  - 2004 – Google releases “MapReduce: Simplified Data Processing on Large Clusters” – large scale processing
  - At the same time – Doug Cutting starts working on a web indexing project called Nutch
    - Inspired by Google papers, he decides to incorporate the storage and processing principles, which are later moved to a new project called Apache Hadoop
  - 2006 – Hadoop is born as an open source project under the Apache Software Foundation

### Big Data and Hadoop History

- At the same time as the Hadoop project was created, other technology innovations were happening (data deluge)
  - Rapid expansion of ecommerce
  - Birth and rapid growth of the mobile Internet
  - Blogs and user-drive web content
  - Social media
- Besides Hadoop, it also led to the emergence of other projects, such as Spark, Messaging Systems, NoSQL, etc

### What is Hadoop?

- Hadoop is a **data storage and processing platform**
- Central concept → **data locality**
  - Data locality refers to the processing of data where it resides (computation is sent to the data, rather than computation requesting data from its location - e.g. DBMS)
- Big data makes it really difficult to move data around
- Hadoop enables processing large datasets locally on the nodes of a cluster
  - Uses a **shared nothing approach**
  - Each node independently processes a subset of the entire dataset without needing to communicate with one another


### What is Hadoop?

- Hadoop is schemaless with respect to its write operations
  - **It is a *schema-on-read* system**
  - Can store and process a wide range of data, such as:
    - Unstructured text documents
    - Semi-structured JSON or XML documents
    - Well structured extracts from relational database systems
- In *schema-on-write* systems, such as relational databases, data is strongly typed and a schema is predefined and enforced in all operations
  - NoSQL databases are *schema-on-read* systems
- Pros: doesn't need indexes, statistics or optimization constructs
- Cons: datasets with larger size and replicated data
- Hadoop uses a **divide and conquer approach** to tackle problems applying data locality and shared nothing concepts





## HDFS Cluster Overview



## YARN Basic Concepts

- Uses a master-slave model like HDFS
  - ResourceManager and NodeManager
- ResourceManager
  - Master of the YARN Cluster
  - Responsible for scheduling applications and granting cluster compute resources
    - Tracks applications' status
  - Resources are represented as units called "containers" (combinations of CPU cores and memory)
    - Tracks available capacity in the cluster - allocates and releases containers to applications
  - For each application submitted, ResourceManager allocates the first container on an available NodeManager as a delegate process for the application called "ApplicationMaster"
  - ApplicationMaster negotiates further containers required for the application
  - ResourceManager serves a web interface on port 8088


## ResourceManager Web Interface

The screenshot shows the 'All Applications' page in the ResourceManager web interface. It displays a table with columns for Application ID, Name, Application Type, Queue, Status, and Progress. The table lists several applications, including 'hadoop-job-1' and 'hadoop-job-2', with their respective statuses and progress bars.

## YARN Basic Concepts


- NodeManager
  - Slave node daemon
  - Manages containers on the host
- ApplicationMaster
  - First container allocated by the ResourceManager to run on a NodeManager for an application
  - Determines resources required to run the application (based on the amount of data to be processed)
  - Orchestrates the allocation of resources to execute application stages (eg. Map and reduce stages)
    - ApplicationMaster requests resources from the ResourceManager on behalf of the application
    - ResourceManager grants resources on the same or other NodeManagers to the ApplicationMaster
  - Monitors progress of tasks, stages (groups of tasks executed in parallel), and dependencies

## YARN Cluster Overview



## Hadoop Cluster Architecture

- Combines HDFS and YARN Clusters



## Hadoop Cluster Deployment Modes

- Hadoop supports three deployment modes
  - LocalJobRunner (Standalone Operation)
    - Runs an application locally in a single JVM
    - Enables integrating the application with an IDE (e.g. Eclipse) for performing unit tests, debugging, and tracing
    - Uses the local filesystem instead of HDFS
  - Pseudo-Distributed
    - All Hadoop daemons are executed on separate JVMs on a single host
    - Simulates how a fully cluster would function
    - Useful for testing to simulate the interaction between components in a real cluster using a single machine
  - Fully Distributed Cluster
    - Executes master and slave daemons on distinct machines
    - Typical deployment mode for production systems

## Pseudo-Distributed Mode



## Hadoop Deployment

- Hadoop was originally developed for Linux
  - Windows-compatible distribution already exists, but may present compatibility issues with other Hadoop ecosystem projects
- Recommendations
  - Do not use LVM (Logical Volume Manager). It may restrict performance (especially on slave nodes)

## Requirements (medium to large scale production Hadoop clusters)


- Master Nodes**
  - 16 or more CPU cores (preferably 32)
  - 128GB or more RAM (preferably 256GB)
  - RAID Hard Drive Configuration (preferably with hot-swappable drives) - for fault tolerance
  - Redundant power supplies - for fault tolerance
  - Bonded Gigabit Ethernet or 10Gigabit Ethernet

## Requirements (medium to large scale production Hadoop clusters)

- Slave Nodes**
  - 16-32 CPU cores
  - 64-512 GB of RAM
  - 12-24 1-4 TB hard disks in a JBOD Configuration
    - JBOD (Just a Bunch of Disks) – directly attached storage not in RAID configuration; each disk operates independently
    - RAID is not recommended for block storage on slave nodes!
      - Access speed is limited by the slowest disk in the array. JBOD has been proven to outperform RAID 0 for block storage by 30% to 50% in benchmarks conducted at Yahoo!
- Slave nodes support failures, therefore they do not require the same degree of fault tolerance that master nodes do
- Storage capacity in slave nodes should not be fully allocated!
  - Failure of a node may impact the network due to replication of blocks

## Networking Considerations

- Fully-distributed implementations consume significant amount of network resources
  - Control messages, status updates and heartbeats, block reports, data shuffling, and block replication
- On-premises deployment
  - Private subnets with dedicated switches
  - Redundant core and TOR (Top of Rack) switches
  - Utilization of STP (Spanning Tree Protocol) to avoid loops
- Hostname resolution between nodes
  - Forward and reverse DNS configuration or hosts file
- Time synchronization (NTP – Network Time Protocol)
  - Used by some components (e.g. Kerberos)



Gerência de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

## Software requirements

- Hadoop is almost entirely written in Java
- Requirements
  - Java Runtime Environment (JRE) 1.7 or above
  - Java Development Kit (JDK) 1.7 or above - for compiling MapReduce applications
- OpenJDK and Oracle versions supported
- Some versions may present problems!
- Site with tested JDKs
  - <https://wiki.apache.org/hadoop/HadoopJavaVersions>
- Other ecosystem projects may have other specific requirements
  - For instance, Apache Spark requires Scala and Python

Gerência de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

## Installing Hadoop On-premises

- Non-Commercial Hadoop
  - Install from Apache builds – <http://hadoop.apache.org>
- Commercial Distributions
  - Cloudera (<http://www.cloudera.com>)
    - Provides a Quickstart VM – <http://www.cloudera.com/downloads.html>
    - Cloudera Manager (proprietary software) – provides a management console and framework to deploy, manage and monitor Hadoop clusters
  - Hortonworks (<http://hortonworks.com>)
    - Hortonworks Sandbox – VM with Hadoop pseudo-distributed mode – <http://hortonworks.com/products/sandbox/>
    - Hortonworks Data Platform (distribution with complete Hadoop stack and selected ecosystem components)
      - Uses Apache Ambari for deployment and management
  - MapR (<https://www.mapr.com/>)
    - MapR Sandbox – VM with Hadoop pseudo-distributed mode – <https://www.mapr.com/products/mapr-sandbox/>
    - MapR Control System (MCS) – central console to configure, monitor and manage MapR clusters

Gerência de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

## Hadoop on the Cloud

- Amazon Web Services
  - EMR (Elastic MapReduce) – Amazon's Hadoop-as-a-Service platform
  - Enables creating clusters with a specific number of nodes, node instance types (EC2), Hadoop distribution, and additional ecosystem applications
  - Uses S3 for reading/writing data
  - Provisioned on demand
  - <https://aws.amazon.com/emr/>

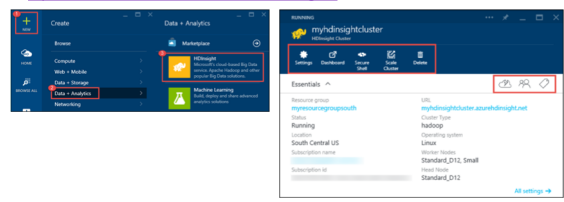


<https://www.youtube.com/watch?v=S6,b05n-c0M>

Gerência de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

## Hadoop on the Cloud

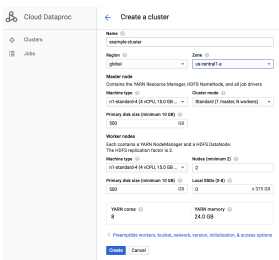
- Microsoft Azure HDInsight
  - <https://azure.microsoft.com/en-us/services/hdinsight/>



Gerência de Infraestrutura para Big Data – Prof. Tiago Ferreira – PUCRS

## Hadoop on the Cloud

- Google Cloud Dataproc
  - <https://cloud.google.com/dataproc/>



# APACHE HADOOP - HANDS-ON