

GERÊNCIA DE INFRAESTRUTURA PARA BIG DATA

Prof. Tiago Ferreto – tiago.ferreto@puccs.br



HDFS – DATA INGESTION

Data Ingestion

- How to capture data produced from source systems in real time?
 - Examples: web logs, databases, sensors
- HDFS standard interface is not practical enough for more complex scenarios
- Tools
 - Flume
 - Sqoop
 - HDFS RESTful interfaces
 - WebHDFS
 - HttpFS


Flume

- Hadoop ecosystem project
 - <https://flume.apache.org/>
- Developed originally by Cloudera
- Goal: capture, transform, and ingest data into HDFS using one or more agents
- Initial use case: capturing log files, or web logs from a web server, and routing them to HDFS as they are generated




Flume Architecture

- Implemented using one or more agents
- Agents connect a data source to HDFS or another agent
- Agents can be chained together or be used in parallel (horizontal scalability or fault tolerance)
- Agents can perform in-flight data operations (e.g., compression, encryption, batching of events, etc)



Flume Agent Architecture

- Flume agents contain a sink, a source, and a channel



Source

- Indicates where the data is to be received from
- Examples
 - HTTP - Used to consume data from RESTful services using POST and GET methods
 - Syslog - Log protocol to capture system events
 - JMS - Java Message Service
 - Kafka - Popular open source messaging platform
 - Avro - Open source, cross platform data serialization framework for Hadoop
 - Twitter - Flume source that connects to Twitter's Streaming API to continuously download tweets

```
agent1.sources = source1
agent1.sources.source1.type = exec
agent1.sources.source1.command = tail -F /tmp/events
```

Sink

- Specifies where to send data (usually HDFS, but can also be another agent or another filesystem, such as S3)
- Examples
 - HDFS - Most common sink used to ingest data into HDFS
 - Hive - Ingests data into HDFS while updating the Hive partitions and the Hive metastore
 - Hbase - NoSQL data store built on HDFS
 - Kafka - Popular open source messaging platform
 - Solr - Open source search platform

```
agent1.sinks = sink1
agent1.sinks.sink1.type = hdfs
agent1.sinks.sink1.hdfs.path = /flume/events
agent1.sinks.sink1.hdfs.filePrefix = events-
agent1.sinks.sink1.hdfs.round = true
agent1.sinks.sink1.hdfs.roundValue = 10
agent1.sinks.sink1.hdfs.roundUnit = minute
```

Channel

- Queue between agent's source and sink
- Flume implements a transactional architecture for reliability (supports rollback and retry operations)
- Possible configurations: **in-memory** or **durable**
- Durable channels use persistent storage (disk) to maintain state (transactional integrity)
- Examples of durable channels
 - File Channel
 - JDBC Channel
 - Kafka Channel

```
agent1.channels = channel1
agent1.channels.channel1.type = memory
agent1.channels.channel1.capacity = 100
agent1.channels.channel1.transactionCapacity = 100
agent1.channels.channel1.byteCapacityBufferPercentage = 20
```

FLUME - HANDS-ON

Sqoop



- Top-level Apache project
 - <http://sqoop.apache.org/>
- Sqoop → sql-to-Hadoop
- Originally developed by Cloudera
- Goal: source data from a relational database and ingest this data into files (typically delimited files) in HDFS
- Can also be used to send data from Hadoop to a relational database
- Integrates with Hive (SQL abstraction to MapReduce)

Sqoop common operations

- Listing databases and tables on a database system
- Importing a single table from a database system, including
 - Specifying which columns to import
 - Specifying which rows to import using a WHERE clause
- Importing data from one or more tables using a SELECT statement
- Incremental imports from a table on a database system (importing only what has changed since a known previous state)
- Exporting of data from HDFS to a table on a remote database system

Gerencia de Infraestrutura para Big Data - Prof. Tiago Ferreira - PUCRS

Sqoop operation

- Implemented using MapReduce (Map-only)
- Steps (import operation)
 1. Connect to the database system using JDBC or a custom connector
 2. Examine the table to be imported
 3. Create a Java class to represent the structure (schema) for the specified table. This class can then be reused for future import operations.
 4. Use YARN to execute a Map-only MapReduce job with a specified number of tasks (mappers) to connect to the database system and import data from the specified table in parallel. The default number of parallel tasks is 4.

The diagram illustrates the Sqoop import process. A client submits a job to the Resource Manager. The Resource Manager then distributes the job into multiple tasks (Task 1, Task 2, Task 3, Task 4) across Hadoop Slave Nodes. Each task independently fetches an equal portion of the dataset from the database and imports it into HDFS. A note indicates that data is imported into HDFS from the parallel portions.

Gerencia de Infraestrutura para Big Data - Prof. Tiago Ferreira - PUCRS

Sqoop commands

Available commands:

<code>codegen</code>	Generate code to interact with database records
<code>create-hive-table</code>	Import a table definition into Hive
<code>eval</code>	Evaluate a SQL statement and display the results
<code>export</code>	Export an HDFS directory to a database table
<code>help</code>	List available commands
<code>import</code>	Import a table from a database to HDFS
<code>import-all-tables</code>	Import tables from a database to HDFS
<code>import-mainframe</code>	Import datasets from a mainframe server to HDFS
<code>job</code>	Work with saved jobs
<code>list-databases</code>	List available databases on a server
<code>list-tables</code>	List available tables in a database
<code>merge</code>	Merge results of incremental imports
<code>metastore</code>	Run a standalone Sqoop metastore
<code>version</code>	Display version information

Gerencia de Infraestrutura para Big Data - Prof. Tiago Ferreira - PUCRS

Sqoop example

- Importing all tables from a MySQL database named "mydb"

```
$ sqoop import-all-tables \
--username javen \
--password ***** \
--connect jdbc:mysql://mydbserver.local/mydb
```

Gerencia de Infraestrutura para Big Data - Prof. Tiago Ferreira - PUCRS

Sqoop2 (Sqoop-as-a-Service)

- Server-based implementation of Sqoop
 - Server hosts the application and connectors, and a lightweight client, web UI, or API connects to the server to submit a Sqoop request
 - Server executes the request on the client's behalf
- Advantages
 - Centralized management of security, as credentials are centrally stored and managed
 - Capability to restrict or control the number of connections to a database system
 - Capability to implement network isolation between client systems and subnets and database systems
 - Eliminates the need to manage software on client systems (e.g., custom connectors, binaries, etc.)

Gerencia de Infraestrutura para Big Data - Prof. Tiago Ferreira - PUCRS

WebHDFS

- WebHDFS provides RESTful access to HDFS using HTTP or HTTPS
- Supports read and write operations
- Command line utilities such as `wget` or `curl` can be used to access HDFS
- Example
 - `$ curl -i -L http://namenode:50070/webhdfs/v1/data/file.txt?op=OPEN`
- Limitations
 - WebHDFS does not support High Availability (HA) HDFS implementations
 - Clients must be able to access every DataNode in the cluster

Gerencia de Infraestrutura para Big Data - Prof. Tiago Ferreira - PUCRS

WEBHDFS - HANDS-ON

HttpFS

- HttpFS provides RESTful access via a service
- Solution is more scalable, supporting HA HDFS implementations and not requiring direct client accessibility to DataNodes in the cluster
- HttpFS server acts as a proxy accepting REST requests from clients and submitting them to HDFS on the clients' behalf



Data Ingestion – Considerations

- When a file is being written to HDFS, the file initially appears to clients as a zero byte file in the incoming directory with a `._COPYING_` suffix
- Data is committed to the file in block size increments (typically 128MB)
 - File size increases in 128MB increments
- The file is visible in the target HDFS directory and can be read (and therefore used as input for processing) while the file is not yet complete
 - Problems can happen processing incomplete files
- Recommendation
 - Use a special directory (e.g., `./incoming`) to ingest data
 - Once the write operation is complete, move the file to an accessible directory (e.g., `./ingested`)
 - Move operation is (almost) instantaneous – only requires a metadata operation