

GERÊNCIA DE INFRAESTRUTURA PARA BIG DATA

Prof. Tiago Ferreto – tiago.ferreto@puccs.br



JAVA MAPREDUCE API

Java MapReduce API – Basic concepts

- **Serialization** is widely present in Java
- **MapReduce API**
 - Conversion of data structures to byte streams and vice versa (deserialization)
- **Hadoop Datatypes**
 - Keys and values are serializable objects
 - Utilization of built-in data object types (Box Classes) instead of primitives types (int, long, char)
 - Requires using specific methods to access/modify data
 - All serializable objects use the **Writable** interface
 - Define accessor and mutator methods
 - Example: methods `readFields` and `write`
 - Sorting is provided by a **WritableComparable** interface
 - Provides methods `compareTo`, `equals` and `hashCode`

Hadoop Box Class	Java Primitive
BooleanWritable	boolean
ByteWritable	byte
IntWritable	int
FloatWritable	float
LongWritable	long
DoubleWritable	double
NullWritable	null
Text	String

Java MapReduce API – Basic concepts

- **InputFormats**
 - Specifies how data (keys and values) are extracted from a file
 - Provide a factory for **RecordReader** objects → used to extract data from an **Input Split** (typically a HDFS block)
- **Common InputFormats**
 - **TextInputFormat**: InputFormat for plain files. Files are broken into lines. Keys are the position in the file, and values are the line of text.
 - **KeyValueTextInputFormat**: Similar to **TextInputFormat**, but with each line divided into a key and value by a separator
 - **SequenceFileInputFormat**: InputFormat for **SequenceFiles** (special Hadoop format)
 - **NLineInputFormat**: Similar to **TextInputFormat**, and specifies how many lines should go to each map task
 - **DBInputFormat**: InputFormat to read data from a JDBC data source
 - **FixedLengthInputFormat**: InputFormat to read input files which contain fixed length records

Java MapReduce API – Basic concepts

- **SequenceFiles**
 - Binary encoded, serialized data files designed for use in Hadoop
 - Contain metadata defining datatypes for key and value objects within the file
 - Can be uncompressed or compressed
 - Efficient in multi-stage workflow (output of one MapReduce used as input for another MR job)
 - Not accessible from other languages (only Java)
 - Alternative: Avro format → provides a cross-language serialization format

Java MapReduce API – Basic concepts

- **OutputFormats**
 - Determine how data is written out to files
- **Common OutputFormats**
 - **FileOutputFormat**: Writes output data to a file
 - **DBOutputFormat**: Writes output data to a JDBC data source

Components of a MapReduce Program

- A typical MapReduce Program contains the following components:
 - Driver – code executed on the client which sets up and starts the MapReduce application
 - Mapper – Java class that contains the map() method
 - Reducer – Java class that contains the reducer() method

Driver

- It submits the application and its configuration to the ResourceManager
- Jobs can be submitted
 - Synchronously: waits for the application to complete before performing another action
 - Asynchronously: does not wait the application to complete
- Can configure and submit more than one application
 - For example: workflow of MapReduce applications
- Typical parameters
 - Path to input data
 - Path to output data
 - HDFS/YARN cluster to be used for the application
- Implemented as a Java class containing the entry point (main() method) for the program

Job Object

- Stores Job configuration
 - Classes to be used as Mapper and Reducer
 - Input and output directories
 - Job name to be displayed in the YARN Resource Manager UI
 - Among other options
- Can also be used to control application's submission and execution and to query the state of the application

Mapper

- Java class containing the map() method
- Each Mapper instance iterates through its assigned InputSplit
 - Executes its map() method against each record read from the InputSplit, using a defined InputFormat and its associated RecordReader
- Number of InputSplits (usually the number of HDFS blocks) in the input data determines the number of Map tasks in a MR application
- Implementation concerns
 - **No saving or sharing of state** – data should not be shared between Map tasks
 - **No side effects** – map tasks may be executed in any sequence or executed more than once without creating side effects
 - **No attempt to communicate with other map tasks** – map tasks are not intended to communicate with one another
 - **Careful on performing external IO operations (writing to NFS or accessing a service)** – can be perceived as a DDoS attack or event storm

Reducer

- Runs against a partition and respective sorted keys
 - Its associated values are passed to the reduce() method
- The same implementation concerns for Mappers should be applied when implementing a reducer

Hello World in MapReduce → Word Count

- Counts the occurrence of words in a text file
- Useful in real-life problems
 - Counting specific event occurrences on log files
 - Text mining functions
 - Word clouds
- Input → text file
- Map task → splits a line of text into a collection of words (tokenization) and outputs each word as key with a value of 1
- Shuffle-and-sort → groups the values for each word and sort keys in alphabetical order
- Reducer → sums the values for each word (key)

```

WordCountDriver.java

public class WordCountDriver extends Configured implements Tool {
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.out.printf("Usage: %s [generic options] <inputdir> <outputdir>\n",
                getClass().getSimpleName()); return -1; }
        Job job = new Job(getConf());
        job.setJarByClass(WordCountDriver.class);
        job.setJobName("Word Count");
        FileInputFormat.setInputPaths(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        job.setMapperClass(WordCountMapper.class);
        job.setReducerClass(WordCountReducer.class);
        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(IntWritable.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        boolean success = job.waitForCompletion(true);
        return success ? 0 : 1;
    }

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new Configuration(), new WordCountDriver(), args);
        System.exit(exitCode);
    }
}

```

```

WordCountMapper.java

public class WordCountMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        String line = value.toString();
        for (String word : line.split("\\W+")) {
            if (word.length() > 0) {
                wordObject.set(word);
                context.write(wordObject, one);
            }
        }
    }
}

```

```

WordCountReducer.java

public class WordCountReducer extends Reducer<Text, IntWritable, Text, IntWritable>
{
    private IntWritable wordCountWritable = new IntWritable();
    @Override
    public void reduce(Text key, Iterable<IntWritable> values, Context context)
        throws IOException, InterruptedException {
        int wordCount = 0;
        for (IntWritable value : values) {
            wordCount += value.get();
        }
        wordCountWritable.set(wordCount);
        context.write(key, wordCountWritable);
    }
}

```

Compiling, Packaging and Submitting

- **Compiling**
 - Put all files in the same folder
 - \$ javac -classpath ` \$HADOOP_HOME/bin/hadoop classpath ` *.java
- **Driver, Mapper and Reducer must be packaged in a jar file**
 - Can use Maven or javac/jar executables
 - \$ jar cvf wc.jar *.class
- **Submission using the hadoop jar command**
 - \$ \$HADOOP_HOME/bin/hadoop jar wc.jar WordCountDriver inputdir outputdir

Advanced MapReduce API Concepts

- **Combiners**
 - Decrease the amount of intermediate data sent between Mappers and Reducers
 - combiner() function is identical to the reduce() function
 - Output key and value object types from map() function match the input to the Combiner
 - Output key and value object types from the Combiner match the input key and value object types used in the Reducer's reduce() method
 - Operation performed must be commutative and associative
- **Declaration in Driver class**
 - job.setCombinerClass(WordCountReducer.class);

Advanced MapReduce API Concepts

- **Partitioners**
 - Divides the output keyspace for a MapReduce application controlling the data each Reducer gets
 - Useful for process distribution, load balancing or segregating output (separating a file for each month of the year)
 - HashPartitioner is used by default
 - Uses a hash function to separate the keyspace in roughly equal parts among Reducers
 - **Declaration of a Custom Partitioner in Driver Class**
 - job.setPartitionerClass(MyCustomPartitioner.class);
 - Extends the base Partitioner class and has a getPartition() method

Advanced MapReduce API Concepts

- **Distributed Cache**
 - Used to disseminate additional data or class libraries at runtime to Mappers or Reducers in a fully distributed Hadoop cluster environment
 - DistributedCache pushes the data or libraries to all slave nodes as a prerequisite background task before any task for the application is executed
 - The distributed data is available in read-only format on any node running a task for the application
 - After the application terminates, the files are removed automatically
 - Items can be added using the ToolRunner class in the Driver (arguments -files, -archives, -libjars)
- **Example: adding stopwords to the WordCount application**

```
$ hadoop jar wc.jar WordCountDriver -files stopwords.txt inputdir outputdir
```
- **Example: accessing files in the DistributedCache**

```
File f = new File("stopwords.txt");
```

MapReduce Streaming API

- Enables implementing Map and Reduce functions in languages other than Java (e.g., Perl, Python, Ruby)
- Requires that key-value pairs are presented to the Map or Reduce script using standard input (STDIN), and key-value pairs are emitted from the script using standard output (STDOUT)
- **Some drawbacks:**
 - Presents additional overhead → performance is poorer than implementing in Java
 - Requires Java for implementing additional API constructs (e.g., InputFormats, Writables, Partitioners)
 - Only suitable for data represented as text

Word Count in Python

```
#!/usr/bin/env python
# wordmapper.py
import sys
for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for word in words:
        print '%s\t%s' % (word, 1)
```

```
#!/usr/bin/env python
# wordreducer.py
import sys
thisword = None
wordcount = 0
word = None
for line in sys.stdin:
    line = line.strip()
    word, count = line.split('\t', 1)
    count = int(count)
    if thisword == word:
        wordcount += count
    else:
        if thisword:
            print '%s\t%s' % (thisword, wordcount)
            wordcount = count
        thisword = word
    if thisword == word:
        print '%s\t%s' % (thisword, wordcount)
```

Submitting a Streaming MapReduce Job

```
$ hadoop jar \
  $HADOOP_HOME/share/hadoop/tools/lib/hadoop-streaming-
*.jar \
  -input inputdir \
  -output outputdir \
  -mapper wordmapper.py \
  -reducer wordreducer.py \
  -file wordmapper.py \
  -file wordreducer.py
```

MAPREDUCE – HANDS-ON