


## GERÊNCIA DE INFRAESTRUTURA PARA BIG DATA

Prof. Tiago Ferreto – tiago.ferreto@puccs.br



### Hive Introduction

- Started in 2010 at Facebook
- Goal: provide a high-level interface to Hadoop MapReduce using SQL-like
  - HiveQL (Hive Query Language) – implements a subset of SQL-92 with extensions
- Process
  - Hive parses the query in HiveQL and generates Java MR operations
  - MR are submitted to the Hadoop cluster
  - Hive monitors and returns results to client (or writes to HDFS)




### Hive Objects

- Hive implements a tabular abstraction to objects in HDFS
  - Directories and files are represented as tables with predefined columns with datatypes
  - Tables are created using SQL DDL (Data Definition Language) commands and accessed using SQL DML (Data Manipulation Language) statements
- Differences between Hive and a conventional database platform
  - UPDATE is implemented as a coarse-grained transformation
    - Since HDFS is an immutable filesystem, modifications to tables are not allowed
  - No transactions, no journaling, no rollbacks, no real transaction isolation level
  - No declarative referential integrity (DRI), no primary keys, no foreign keys
  - Incorrectly formatted data are represented to the client as NULL

### Hive Metastore

- Mapping of tables to directory locations in HDFS, columns and definitions are maintained in Hive metastore
  - Hive Metastore is a relational database accessed by the Hive client
- Object definitions include input/output formats of files (tables) and serialization/deserialization functions to define how records and fields are extracted
- Hive Metastore is usually implemented as an embedded Derby database
  - But can also be implemented as a local or remote database (MySQL, PostgreSQL, etc)
- HCatalog - Hive subproject
  - Extend objects created in Hive to other projects (e.g., Apache Pig)



### Hive CLI

- Accepts and parses HiveQL input commands
- Common method for performing ad hoc queries
- Used when the Hive client or driver application is deployed to the local machine, including the connection to the metastore

```
[java@hadoop-01 ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf/dis
hive> show tables in movielens;
+-----+
| data |
| genre |
| info |
| item |
| occupation |
| user |
+-----+
Time taken: 0.678 seconds, Fetched: 6 row(s)
hive>
```

## HiveServer2

- Client/server approach to use Hive
- Used in large-scale implementations
  - Connection details to the metastore are kept in the server and cluster access is controlled
- Can act as a multi-session driver application for multiple clients
- Provides a JDBC interface to be accessed by external clients (e.g., Beeline – lightweight CLI)

## Hive Batch Mode

- Hive also supports non-interactive or batch mode execution
- Example
  - `# run all statements in a file`
  - `$ hive -f MyHiveQuery.hql`
  - `# run an individual statement`
  - `$ hive -e "SELECT * FROM mytable"`

## Hive Database and Tables

- Hive objects consist basically of databases and Tables
- Databases – used for organization, authorization, and namespace management
  - Default database is named "default"
- Tables – exist in a Hive database
- Database context can be changed using the USE statement
  - `hive> USE bikeshare;`
- Objects can be referenced using fully qualified object identifier (database and table name)
  - `hive> SELECT * FROM bikeshare.trips;`

## Hive Authorization

- Hive supports basic authorization
  - Determines access levels to objects within a Hive instance
  - Not enabled by default
  - Configured in hive-site.xml file

```

<!--
<property>
  <name>hive.security.authorization.enabled</name>
  <value>true</value>
</property>
<property>
  <name>hive.security.authorization.createable.owner.grants</name>
  <value>ALL</value>
</property>
-->

```

- Object-level privileges are assigned using the GRANT statement
  - Operations: SELECT, CREATE, ALTER, DROP, ALL, ...
  - Scope: GROUP, ROLE, USER
  - Revoking privileges: REVOKE statement
  - Examples:
    - `hive> GRANT SELECT ON TABLE trips TO USER javen;`
    - `hive> REVOKE SELECT ON TABLE trips FROM USER javen;`

## Hive Objects

- Hive supports most common primitive datatypes, and also complex datatypes (struct, map, array)

Datatype	Category	Description
TINYINT	Numeric	1-byte signed integer (-128 to 127)
SMALLINT	Numeric	2-byte signed integer (-32,768 to 32,767)
INT	Numeric	4-byte signed integer
BIGINT	Numeric	8-byte signed integer
FLOAT	Numeric	4-byte single precision floating point
DOUBLE	Numeric	8-byte double precision floating point
DECIMAL(p, s)	Numeric	User definable precision and scale
TIMESTAMP	Date/Time	Unix timestamp
DATE	Date/Time	Date (formatted as YYYY-MM-DD)
STRING	String	Character sequence (variable length)
VARCHAR	String	Character sequence (variable length)
CHAR(n)	String	Character sequence (fixed length)
BOOLEAN	Misc	True or False
BINARY	Misc	Raw binary data

## CREATE TABLE Statement

```

hive> CREATE TABLE stations
> (
>   station_id INT,
>   name STRING,
>   lat DOUBLE,
>   longitude INT,
>   dockcount INT,
>   landmark STRING,
>   installation STRING
> )
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS TEXTFILE;

```

```

hive> DESCRIBE stations;
OK
station_id      int
name            string
lat             double
longitude       double
dockcount       int
landmark        string
installation     string
Time taken: 0.035 seconds, Fetched: 7 row(s)

```

## Input/Output Formats and SerDes in Hive

- Hive uses an InputFormat and a SerDe (Serializer/Deserializer) to determine how to read input files and extract records for processing
  - Specified using the STORED AS directive in the CREATE TABLE statement
  - Example: STORED AS TEXTFILE corresponds to the TextInputFormat
- Other InputFormats: SEQUENCEFILE, ORC, PARQUET, AVRO
- Default SerDe used by Hive – LazySimpleSerDe
  - File is delimited by a character (printable or non-printable) indicated in the FIELDS TERMINATED BY clause (defined as Ctrl-A if not present)
- Other SerDe (specified using the ROW FORMAT SERDE directive)
  - RegexSerDe – schema for fixed-width data
- Default OutputFormat → HiveIgnoreKeyTextOutputFormat (used to write data to text files in Hadoop)
- Other OutputFormats: ORC, Parquet, Avro, SequenceFiles, ...

## Loading data into Hive

- Approaches to load data into tables
  - Put, move, or copy the file(s) into the HDFS directory for the table object
  - Use the Hive LOAD DATA INPATH command for data that exists in HDFS, which uses an underlying HDFS move function
  - Use the Hive LOAD DATA LOCAL INPATH command for data that does NOT EXIST in HDFS, which uses an underlying HDFS put function
- Examples
  - hive> **LOAD DATA INPATH** '/bikeshare/stations' INTO TABLE stations;
  - Use **OVERWRITE** to replace exiting contents
    - hive> **LOAD DATA INPATH** '/bikeshare/stations' **> OVERWRITE** INTO TABLE stations;
  - Data can be loaded during table creation process
    - hive> **CREATE TABLE** stations\_copy AS **> SELECT \*** FROM stations;

## Analyzing data with Hive

- Hive Query Language (HiveQL) is based on the SQL-92 specification, with some additional Hive-specific functions
- HiveQL statements can span multiple lines and are terminated by a semicolon.
- Single line comments are supported using the double hyphen (--).
- Typical SQL semantics such as column lists and WHERE clauses are fully supported in Hive
- Example
 

```
-- this is a comment
hive> SELECT name, lat, long
> FROM stations
> WHERE landmark = 'San Jose';
```

## Hive's Built-in Functions

- Hive includes several built-in functions to perform several operations (mathematical, string and date manipulation, etc)
  - Examples: ROUND, CEIL, FLOOR, RAND, SUBSTRING, LOWER, UPPER, RTRIM, LTRIM, CONCAT, TO\_DATE, DAY, MONTH, YEAR
- Many functions are identical to the ones available in the SQL dialects in most popular RDBMS
- If a function is missing, additional UDFs (User Defined Functions) can be written (in Java)
- Examples
  - hive> **SELECT LOWER(name)** FROM stations;
  - DESCRIBE** can be used to show help and usage for functions
  - hive> **DESCRIBE FUNCTION RAND**;
  - OK
  - RAND([seed])** - Returns a pseudorandom number between 0 and 1

## Data grouping and aggregation

```
hive> SELECT landmark, COUNT(*) FROM stations
> GROUP BY landmark;
OK
Mountain View      7
Palo Alto          5
Redwood City       7
San Francisco      35
San Jose           16
Time taken: 19.443 seconds, Fetched: 5 row(s)
```

## Joins in Hive

- Hive supports all common join types (INNER JOIN – default, LEFT OUTER JOIN, RIGHT OUTER JOIN and FULL OUTER JOIN)

```
hive> SELECT t.trip_id,
> t.duration, t.start_date, s.name
> FROM stations s
> JOIN trips t ON s.station_id = t.start_terminal;
```

## Data Output with Hive

- Hive supports several methods to persist the output of a query to local or HDFS filesystem
  - INSERT OVERWRITE** → sends the query results to another Hive table (a directory in HDFS), overwriting the existing contents
  - INSERT INTO TABLE** → sends the query results to another Hive table, appending the output to the existing table (directory) contents
  - INSERT OVERWRITE DIRECTORY** → saves the results to a directory in HDFS which may or may not be assigned to a Hive table
  - INSERT OVERWRITE LOCAL DIRECTORY** → saves the results to a local directory, not to HDFS

```
hive> INSERT INTO TABLE trips_counts
> SELECT start_terminal, start_station, COUNT(*) AS count
> FROM trips
> GROUP BY start_terminal, start_station
> ORDER BY count
> DESC LIMIT 10;
```

## Complex Datatypes in Hive

- Hive supports complex or nested datatypes
  - Common in XML, JSON, and other semi-structured data sources
- ARRAY Datatype**
  - Stores an ordered list of elements of the same datatype
  - Uses **COLLECTION ITEMS TERMINATED BY** clause to determine the delimiter used for reading or deriving elements in the collection

```
hive> CREATE TABLE customers (            hive> SELECT orderids[0] FROM customers;
> id INT,                                1
> fname STRING,
> lname STRING,
> email STRING,
> orderids ARRAY<INT>)
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY '|'
> COLLECTION ITEMS TERMINATED BY ','
> ;
```

## Complex Datatypes in Hive

- STRUCT Datatype**
  - Consists of named fields that can be of the same of different datatypes
  - Uses angle brackets (<>) to associate datatypes to fields

```
hive> CREATE TABLE customers (
> id INT,
> fname STRING,
> lname STRING,
> email STRING,
> orderids ARRAY<INT>,
> email_preferences STRUCT<opcomms:boolean, promos:boolean>
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY '|'
> COLLECTION ITEMS TERMINATED BY ','
> ;
hive> SELECT email_preferences.opcomms FROM customers;
TRUE
```

## Complex Datatypes in Hive

- MAP Datatype**
  - Represents key value pairs

```
hive> CREATE TABLE customers (
> id INT,
> fname STRING,
> lname STRING,
> email STRING,
> orderids ARRAY<INT>,
> email_preferences STRUCT<opcomms:boolean, promos:boolean>,
> address_map MAP<STRING, STRING>
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY '|'
> COLLECTION ITEMS TERMINATED BY ','
> MAP KEYS TERMINATED BY ':'
hive> SELECT address_map['city'] FROM customers;
Hayward
```


## Query management and optimization in Hive

- Hive requires optimization in queries design to improve performance
  - Unlike a traditional relational database platform, Hive does not have indexes and statistics to optimize queries and data access
- An important factor for optimization is the amount of data read by a particular query
- Hive offers two key approaches used to limit or restrict the amount of data that a query needs to read
  - Partitions
  - Buckets

## Partitions in Hive

- Used to divide data into subdirectories based upon one or more conditions (typically used in **WHERE** clauses)
- Used normally for coarse-grained data grouping
- Example**


```
hive> CREATE TABLE weblogs (
> ip STRING,
> time_local STRING,
> method STRING,
> uri STRING,
> protocol STRING,
> status STRING,
> bytes_sent STRING,
> referer STRING,
> useragent STRING)
> PARTITIONED BY (date STRING);
```



## Buckets in Hive

- Used to group items based upon a key hash

```
hive> CREATE TABLE customers
> (cust_id INT,
> fname STRING,
> lname STRING,
> email STRING
> )
> CLUSTERED BY (cust_id) INTO 3 BUCKETS;
```



## EXPLAIN Statement

- Used to understand the query execution plan of the Hive query optimizer for a specific query

```
hive> EXPLAIN SELECT method, COUNT(*) FROM weblogs GROUP BY method;
STAGE DEPENDENCIES:
  Stage-1 is a root stage
  Stage-0 depends on stages: Stage-1
STAGE PLANS:
  Stage: Stage-1
    Map Reduce
    Map Operator Tree:
      TableScan
      alias: weblogs
      ...
```

HIVE – HANDS-ON